

# HW4 — ISYE6416

Ashish Dhiman — ashish.dhiman9@gatech.edu

May 7, 2023

## Contents

<b>1</b>	<b>Question 1: Conceptual questions</b>	<b>2</b>
1.1	Data fit complexity in regression tree . . . . .	2
1.2	Bagging and Boosting . . . . .	2
1.3	OOB Error . . . . .	2
1.4	Optimal $\alpha_t$ . . . . .	2
<b>2</b>	<b>Question 2: Random forest for email spam classifier</b>	<b>4</b>
2.1	CART model . . . . .	4
2.2	Random Forest . . . . .	4
2.3	Random Forest . . . . .	5
<b>3</b>	<b>Question 3: AdaBoost</b>	<b>6</b>
3.1	Adaboost iterations fitting . . . . .	6
3.2	Final classifier . . . . .	9
<b>4</b>	<b>Question 4: Importance Sampling</b>	<b>10</b>
4.1	Importance Ratio . . . . .	10
4.2	Importance Sampling Algorithm . . . . .	10
4.3	Numerical Example . . . . .	11
<b>5</b>	<b>References:</b>	<b>11</b>

# 1 Question 1: Conceptual questions

## 1.1 Data fit complexity in regression tree

We control the data fit complexity in a regression tree by first growing the tree greedily until a minimum node size is reached. Then pruning of the tree is performed by optimising the cost function:

$$C(\alpha) = MSELoss(J) + \alpha(|nodes_j|)$$

where  $nodes_j$  is number of terminal nodes in sub tree J.

We can control overfitting by various parameters such as:

1. cost complexity parameter or  $\alpha$ , increasing  $\alpha$  decreases overfitting
2. min number of children of a node, again this limits the depth of the tree and only allows further splits, if there is enough population at a node.

The data fit complexity in a regression tree can be controlled by a variety of parameters such as:

## 1.2 Bagging and Boosting

Bagging and Boosting are both ensemble techniques, where multiple models are built and then combined to get prediction.

While in Bagging the different models are independent and the end predictions are simple averaged.

On the other hand in case of Boosting the individual models are not independent, and the successive models put more weight on erroneous data points from previous models. The result at the end too is a weighted average of the individual model predictions.

## 1.3 OOB Error

While building Random forest, each individual tree model is built on data sampled randomly with replacement, or bootstrapped data. The data points which do not get sampled for this particular model are referred to as Out of Bag points, or OOB.

When we evaluate the individual model on these OOB points we land up with OOB error rate. Because the OOB points were not used to train the original model, OOB error is similar to the Test error rate.

We can thus see the drop OOB error with increasing number of trees, and choose the point where the drop in OOB error vanishes.

## 1.4 Optimal $\alpha_t$

We know

Classifier at iteration t is :

$$f_t(x) = \sum_{i=1}^t \alpha_i h_i(x)$$

Exponential Loss for above classifier is:

$$\begin{aligned} L_t(f) &= \frac{1}{m} \sum_{i=1}^m (\exp(-y_i * f_t(x_i))) \\ L_t(f) &= \frac{1}{m} \sum_{i=1}^m (\exp(-y_i * [f_{t-1}(x_i) + \alpha_t h_t(x_i)])) \\ &= \frac{1}{m} \sum_{i=1}^m [(\exp(-y_i * f_{t-1}(x_i)) * \exp(-y_i * \alpha_t h_t(x_i)))] \end{aligned}$$

Defining weight for i'th point at iteration t as :

$$\begin{aligned} D_t(i) &= \frac{\exp(-y_i * f_{t-1}(x_i))}{Z_t} \\ \implies L_t(f) &= \frac{1}{m} \sum_{i=1}^m [(D_t(i) * Z_t * \exp(-y_i * \alpha_t h_t(x_i)))] \end{aligned}$$

Taking derivative wrt  $\alpha$

$$\begin{aligned} \frac{\partial L}{\partial \alpha_t} &= \frac{-1}{m} \sum_{i=1}^m [(D_t(i) * Z_t * \exp(-y_i * \alpha_t h_t(x_i)) * y_i * h_t(x_i))] \\ &= \frac{-1}{m} \sum_{y_i=h_i} (D_t(i) * \exp(-\alpha_t)) + \frac{1}{m} \sum_{y_i \neq h_i} (D_t(i) * \exp(\alpha_t)) \end{aligned} \tag{1}$$

We also have:

$$\epsilon_t = \sum_{i=1}^m D_t(i) \mathcal{I}(y_i \neq h_t(x_i))$$

Therefore :

$$\begin{aligned} \implies \epsilon_t &= \sum_{y_i \neq h_i} D_t(i) \\ \implies 1 - \epsilon_t &= \sum_{y_i = h_i} D_t(i) \\ &= \frac{(\epsilon_t - 1) * \exp(-\alpha_t)}{m} + \frac{(\epsilon_t) * \exp(\alpha_t)}{m} \\ \implies \frac{1 - \epsilon_t}{\epsilon_t} &= \exp(2\alpha_t) \\ \implies \alpha_t &= \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t} \end{aligned}$$

## 2 Question 2: Random forest for email spam classifier

### 2.1 CART model

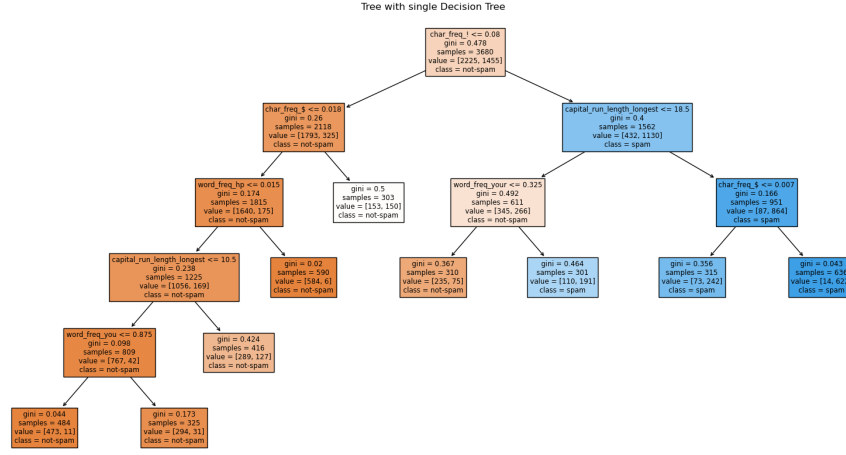


Figure 1: Tree Model with single Decision tree

### 2.2 Random Forest

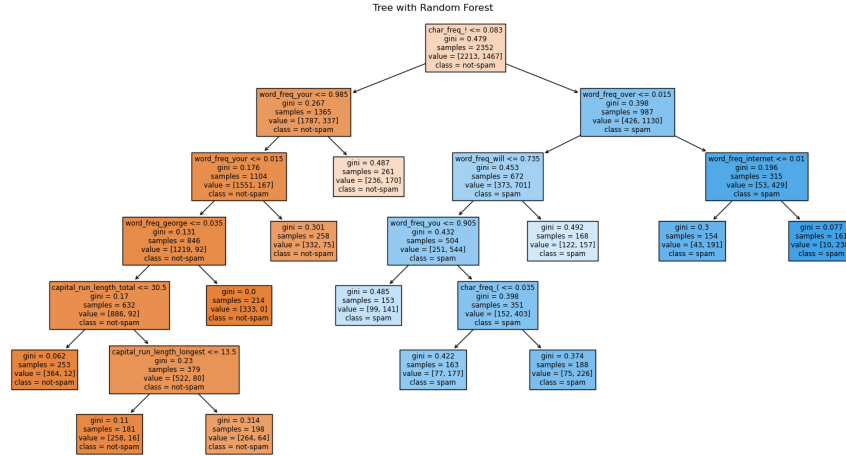


Figure 2: Tree Model with Random Forest

A good rule of thumb value is:

- For Classification,  $\nu \approx \sqrt{p}$
- For Regression,  $\nu \approx \lfloor p \rfloor$

## 2.3 Random Forest

With 80% data for Train and 20% data for Test the accuracy results are as follows:

- Decision Tree, 80.8%
- Random Forest, 92.1%

The Accuracy is better with Random Forest, and both the values can be improved further with proper tuning of model hyper parameters.

The OOB error curve is given below:

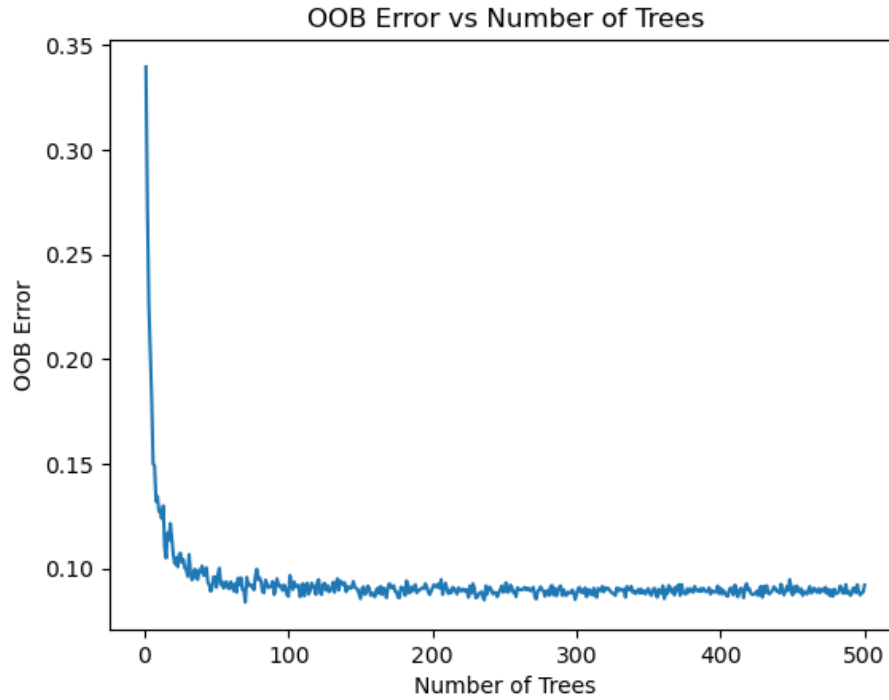


Figure 3: OOB error curve

## 3 Question 3: AdaBoost

### 3.1 Adaboost iterations fitting

I am using the following code to fit and calculate Adaboost params:

```
1 def get_eps(D,y,h):
2     ind_yh = (y!=h).astype("int")
3     return (np.sum(D*ind_yh))
4
5 def get_alpha(eps):
6     return 0.5*np.log((1-eps)/(eps))
```

#### Iteration 1

```
1 h1 = np.array([1,1,-1,-1,-1,-1,-1,-1])
2 i = 1
3
4 get_base_fig()
5 plt.axvline(-0.25, color="black",lw = 4)
6 plt.title(f"iteration {i}")
7
8 w1 = np.array([1]*8)
9 w1 = w1/w1.sum()
10 eps1 = get_eps(w1,y,h1)
11 alpha1 = get_alpha(eps1)
12
13 print(f"$\epsilon_{i}$ = {eps1:.5f}$")
14 print(f"$D_{i}$: {w1}$")
15 print(rf'$\alpha_{i}$ = {alpha1:.5f}$')
```

$$\epsilon_1 = 0.25000$$

$$D_1 : [0.1250.1250.1250.1250.1250.1250.125]$$

$$\alpha_1 = 0.54931$$

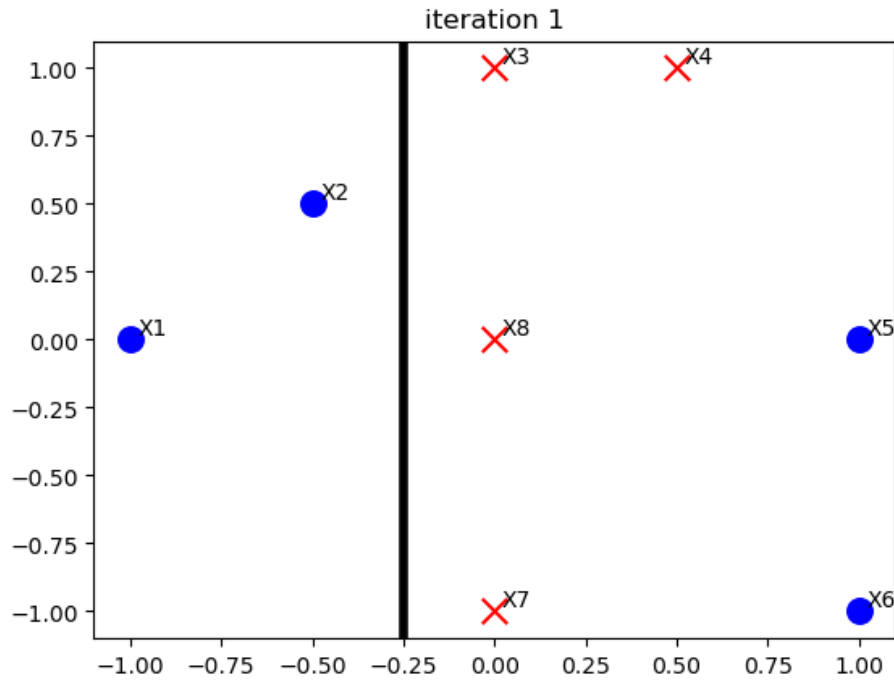


Figure 4: Decision Stump for iteration 1

## Iteration 2

```

1 h2 = np.array([-1,-1,-1,-1,1,1,-1,-1])
2 i = 2
3
4 get_base_fig()
5 plt.axvline(0.75, color="black",lw = 4)
6 plt.title(f"iteration {i}")
7
8 w2 = w1*np.exp(-alpha1*h1*y)
9 z1 = w2.sum()
10 w2 = w2/z1
11 eps2 = get_eps(w2,y,h2)
12 alpha2 = get_alpha(eps2)
13
14 print(f"\epsilon_{i} = {eps2:.5f}$")
15 print(f"$D_{i}$: {w2}$")
16 print(f"$Z_{i-1}$: {z1}$")
17 print(rf'$ \alpha_{i} = {alpha2:.5f}$')

```

$$\epsilon_2 = 0.16667$$

$$D_2 : [0.083333330.083333330.083333330.083333330.250.250.083333330.08333333]$$

$$Z_1 : 0.86603$$

$$\alpha_2 = 0.80472$$

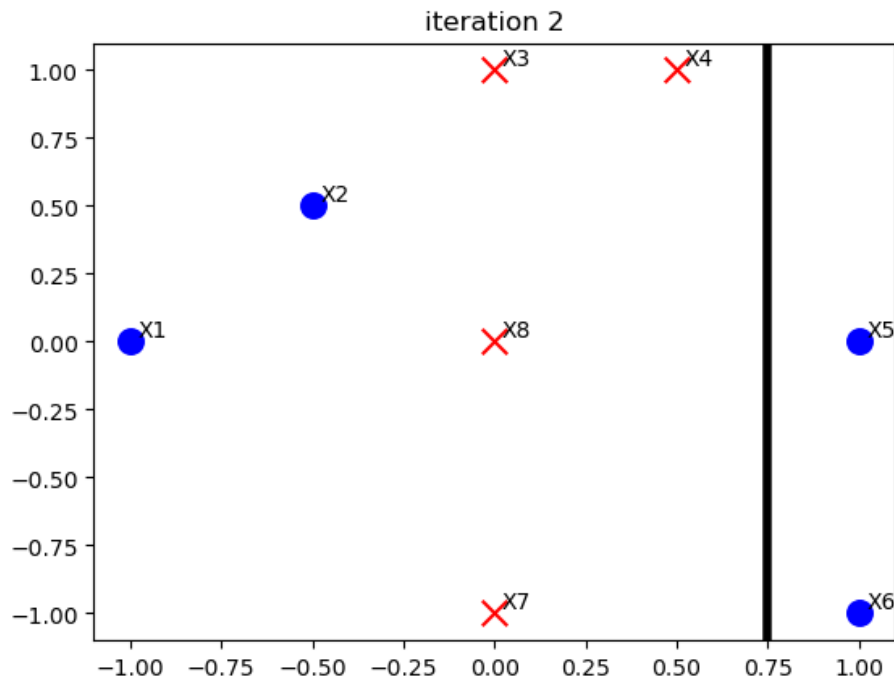


Figure 5: Decision Stump for iteration 2

### Iteration 3

```

1 h3 = np.array([1,1,-1,-1,1,1,1,1])
2 i = 3
3
4 get_base_fig()
5 plt.axhline(0.75, color="black",lw = 4)
6 plt.title(f"iteration {i}")
7
8 w3 = w2*np.exp(-alpha2*h2*y)
9 z2=w3.sum()
10 w3 = w3/z2
11 eps3 = get_eps(w3,y,h3)
12 alpha3 = get_alpha(eps3)
13
14 print(f"\epsilon_{i} = {eps3:.5f}$")
15 print(f"$D_{i}$: {w3}$")
16 print(f"$Z_{i-1}$: {z2:.5f}$")
17 print(rf'$\alpha_{i} = {alpha3:.5f}$')

```

$\epsilon_3 = 0.10000$   
 $D_3 : [0.250.250.050.050.150.150.050.05]$   
 $Z_2 : 0.74536$   
 $\alpha_3 = 1.09861$

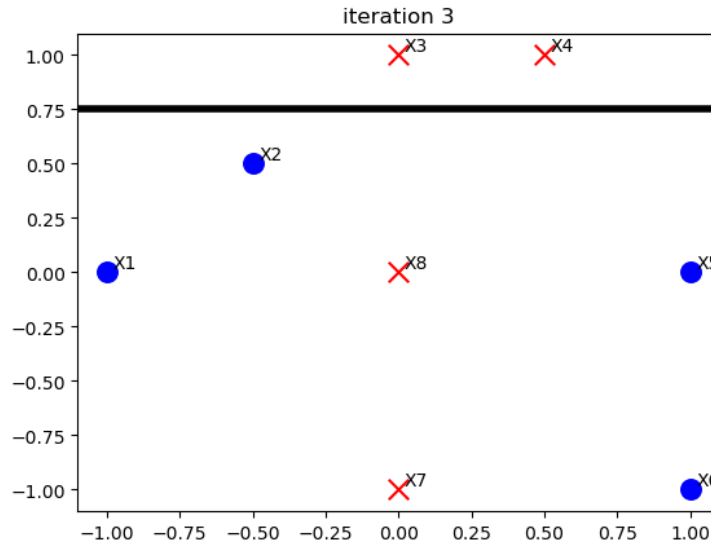


Figure 6: Decision Stump for iteration 3

Similarly  $Z_3 = 0.59999$

$i$	$\epsilon_t$	$\alpha_t$	$Z_t$	$D_t(1)$	$D_t(2)$	$D_t(3)$	$D_t(4)$	$D_t(5)$	$D_t(6)$	$D_t(7)$	$D_t(8)$
1	0.25	0.54931	0.86603	0.125	0.125	0.125	0.125	0.125	0.125	0.125	0.125
2	0.16667	0.80472	0.74536	0.0833	0.0833	0.0833	0.0833	0.25	0.25	0.0833	0.0833
3	0.1	1.09861	0.6	0.25	0.25	0.05	0.05	0.15	0.15	0.05	0.05

Table 1: AdaBoost iteration results



### 3.2 Final classifier

The final classifier is  $H_f = \text{Sign}(\sum_{t=1}^3 \alpha_t * h_t)$

$$H_f = 0.54931 * h_1 + 0.80472 * h_2 + 1.09861 * h_3$$

with

$$h_1 = [1, 1, -1, -1, -1, -1, -1, -1]$$

$$h_2 = [-1, -1, -1, -1, 1, 1, -1, -1]$$

$$h_3 = [1, 1, -1, -1, 1, 1, 1, 1]$$

Hence we get

$$H_f = [1., 1., -1., -1., 1., 1., -1., -1.]$$

Thus our classification accuracy is 100%.

Note: This accuracy is on train data.

Adaboost is a common ensemble learning method that combines numerous weak learners to generate a strong learner. Because Adaboost combines many weak learners, each of which might make mistakes on separate areas of the dataset it is able to increase overall accuracy by combining these learners.

The individual decision trees are also prone to be influenced by outliers, however the combination of multiple learners with different data points weighted differently is unlikely to do so.

## 4 Question 4: Importance Sampling

We have a RV  $X$  distributed as  $\mathcal{N}(\mu_0, 1)$ . We want the right tail probability  $\alpha = \mathbb{P}\{X \geq z\}$ .

For  $z \gg \mu_0$ ,  $\alpha$  is very small, and estimating this small probability accurately is not easy. Now to improve accuracy, we sample from another Gaussian random variable  $Y$  with mean  $\mu_1 = z$  and variance equal to 1.

### 4.1 Importance Ratio

Importance Ratio is given by the ratio of likelihood of saying values above  $z$  under the original distribution  $\mathcal{N}(\mu_0, 1)$  vs the new distribution  $\mathcal{N}(z, 1)$ .

$$\begin{aligned} \text{We have } X &\sim \mathcal{N}(\mu_0, 1), Y \sim \mathcal{N}(z, 1) \\ \implies \text{importance}_{ratio} &= \frac{P(X \geq z)}{P(Y \geq z)} \\ &= \frac{1 - P(X \leq z)}{1 - P(Y \leq z)} \\ \implies \text{importance}_{ratio} &= \frac{1 - F_x(z)}{1 - F_y(z)} \end{aligned} \tag{2}$$

where  $F_x, F_y$  are CDF for  $X, Y$  respectively

### 4.2 Importance Sampling Algorithm

The algorithm is given below:

---

**Algorithm 1** Importance Sampling

---

**Require:**  $\mu_0, z, N$

Initialise Y-list=[], importance-ratio-list = []

**for**  $i = 1$  to  $N$  **do**

$y_i \leftarrow \mathcal{N}(z, 1)$  (Sample from  $Y$ )

$\text{importance}_{ratio}(i) \leftarrow \frac{1 - F_x(z)}{1 - F_y(z)}$  (calculate importance ratio)

    Append  $y_i$  and importance ratio to lists initialised above

**end for**

$I \leftarrow \frac{1}{N} \sum_i NI(y_i \geq z) * \text{importance}_{ratio}(i)$

---

The python implementation is as follows:

```
1 def importance_sampling(mu_x, z, N=100):
2     x_f = norm(loc=mu_x, scale = 1)
3     y_g = norm(loc=z, scale = 1)
4     # draw monte carlo samples
5     x_samples = x_f.rvs(size = N)
6     y_samples = y_g.rvs(size = N)
7     #get importance ratio
8     pr_f_x = x_f.sf(x=y_samples) #Pr(x>=z) from original
9     pr_g_y = y_g.sf(x=y_samples) #Pr(x>=z) from sampling
10    importance_rt = pr_f_x/pr_g_y
11    #approximate I1
12    indicataor_x = (x_samples>=z).astype("int")
13    I1 = np.mean(indicataor_x)
14    #approximate I2
15    indicataor_y = (y_samples>=z).astype("int")
16    I2 = np.mean(indicataor_y*importance_rt)
17    return I1, I2
```

### 4.3 Numerical Example

For  $\mu_0 = 1$  and  $z = 10.0$ , we get the following results:

```
1 True value 1.1285884059538324e-19
2 For Monte Carlo, mean = 0.0, std = 0.0
3 For Importance Sampling, mean = 1.0772267698651663e-20, std =
  3.469453274231522e-21
```

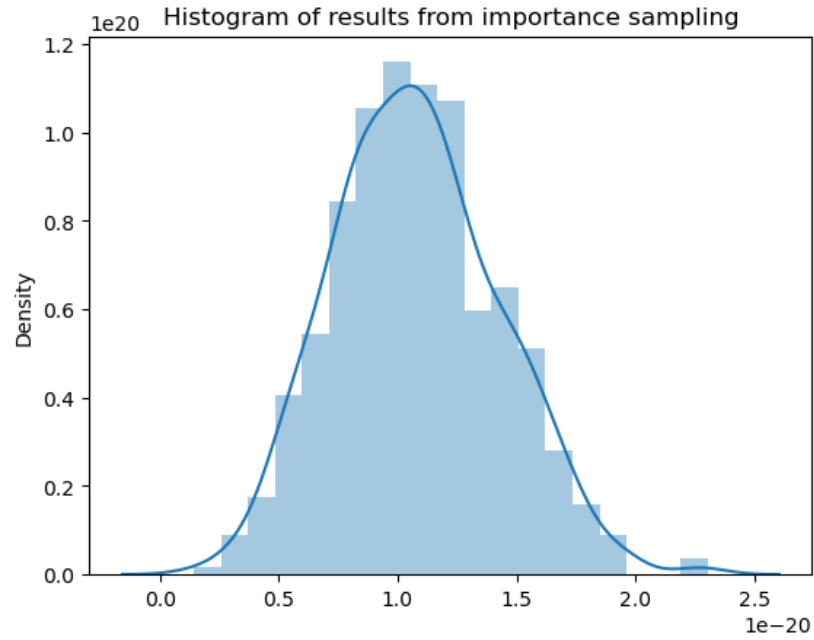


Figure 7: Distribution of results from importance sampling

## 5 References:

1. **Collaborators:** Yibei, Dipendra